1 Understanding Vector Increment

Listing 1: incr.cpp

```
#include <cstring>
#include <cstdio>
#include <cstdlib>
#include <string>

#include "clhelp.h"

int main(int argc, char *argv[])
{
  std::string incr_kernel_str;

  /* Provide names of the OpenCL kernels
   * and cl file that they're kept in */
  std::string incr_name_str =
    std::string("incr");
  std::string incr_kernel_file =
    std::string("incr.cl");

  cl_vars_t cv;
  cl_kernel incr;

  /* Read OpenCL file into STL string */
  readFile(incr_kernel_file,
           incr_kernel_str);

  /* Initialize the OpenCL runtime
   * Source in clhelp.cpp */
  initialize_ocl(cv);

  /* Compile all OpenCL kernels */
  compile_ocl_program(incr, cv, incr_kernel_str.c_str(),
                      incr_name_str.c_str());

  /* Arrays on the host (CPU) */
  float *h_Y, *h_YY;
  /* Arrays on the device (GPU) */
  cl_mem g_Y;

  int n = (1<<20);
  h_Y = new float[n];
  h_YY = new float[n];

  for(int i = 0; i < n; i++)
    {
      h_YY[i] = h_Y[i] = (float)drand48();
    }

  cl_int err = CL_SUCCESS;
  /* CS194: Allocate memory for arrays on
```

```
 * the GPU */
/* Creates a buffer in the cv.context context, with read and write
   access
 * at the global host adress g_Y, of size sizeof(float)*n. */
g_Y = clCreateBuffer(cv.context,CL_MEM_READ_WRITE,sizeof(float)*n,NULL,&
   err);
CHK_ERR(err);

/* enqueue commands to write to the buffer g_Y from hos memory.
 * Commands will be queued in cv.commands.
 * true indicates that the write is put on the commands queue.
 * 0 is the offset in bytes in the buffer object to write to.
 * sizeof(float)*n is the size in byte of data being wirtten.
 * h_Y is the address in host memory of the data being written from.
 */
 err = clEnqueueWriteBuffer(cv.commands, g_Y, true, 0, sizeof(float)*n,
                            h_Y, 0, NULL, NULL);
 /* checks whether the write buffer command was succesful. */
CHK_ERR(err);

/* declaring the global size of th y dimension to be n. */
size_t global_work_size[1] = {n};
/* declaring the size of work groups to be 128 work items. */
size_t local_work_size[1] = {128};

/* Sets specific arguments for the kernel incr.
 * 0 is the argument index, sizeof(cl_mem) is the size
 * of the argument, which is the pointer to g_Y.*/
err = clSetKernelArg(incr, 0, sizeof(cl_mem), &g_Y);
CHK_ERR(err);

/* Sets specific arguments for the kernel incr.
 * 1 is the argument index, sizeof(int) is the size
 * of the argument, which is the pointer to n.*/
err = clSetKernelArg(incr, 1, sizeof(int), &n);
CHK_ERR(err);

/* Enqueues a command on cv.commands to execute the
 * kernel incr.cl on the device. Uses linear dimension
 * to specify work groups and items and specifies to use
 * global_work_size work items for the execution and local_work_size
 * as the size of a work group.  */
err = clEnqueueNDRangeKernel(cv.commands,
                             incr,
                             1,//work_dim,
                             NULL, //global_work_offset
                             global_work_size, //global_work_size
                             local_work_size, //local_work_size
                             0, //num_events_in_wait_list
                             NULL, //event_wait_list
                             NULL //
                             );
CHK_ERR(err);
```

```
  /* Read result of GPU on host CPU */
  err = clEnqueueReadBuffer(cv.commands, g_Y, true, 0, sizeof(float)*n,
                            h_Y, 0, NULL, NULL);
  CHK_ERR(err);

  /* Check answer */
  bool er = false;
  for(int i = 0; i < n; i++)
    {
      float d = (h_YY[i] + 1.0f);
      if(h_Y[i] != d)
        {
          printf("error at %d :(\n", i);
          er = true;
          break;
        }
    }
  if(!er)
    {
      printf("CPU and GPU results match\n");
    }

  uninitialize_ocl(cv);

  delete [] h_Y;
  delete [] h_YY;

  clReleaseMemObject(g_Y);
  return 0;
}
```

Listing 2: incr.cl

```
/* The __kernel qualifier declares a function
 * that can be executed by an application running
 * on an OpenCL device.
 * The __global qualifier declares that the pointer
 * to Y can point only to the global memory pool.
 * i.e. Y must be in the global memory pool.*/
__kernel void incr (__global float *Y, int n)
{
  /* get_global_id(0) returns the global index of the
   * of the current work item. The 0 argument indicates
   * dimension 0. You can give dimensional indices to
   * work items. In this case it is linear. */
  int idx = get_global_id(0);
  /* If the global index of the work item is less than
   * the size of the array at Y, add 1 to the value of the
   * at adress Y + idx. */
  if(idx < n)
    {
      Y[idx] = Y[idx] + 1.0f;
    }
}
```